

Reduce memory consumption during content-build #15601

🔗 Merged

ncksllyn merged 8 commits into master from nps--fix-memory-issue 📄 on Jan 13


🏠 Code Climate ✓

Conversation 24

Commits 8

Checks 4

Files changed 24

 ncksllyn commented on Jan 11 • edited ▼

Contributor

Description

[department-of-veterans-affairs/va.gov-team#13821](#)

The big memory hog during our build is when we parse every single HTML file into a DOM using Cheerio and store the entire result in memory at once. This PR modifies that by parsing each DOM one at a time, doing the next build steps that require the DOM, then immediately tearing down the DOM so that we only ever have one DOM tree in memory at once. I had to refactor the build steps (those dependent on `file.dom`) to accept just a single `file` at a time to make this work.

It appears we can save several gigs of memory and shave off almost 30 seconds doing this.

▼ Build output w/these changes

Step Step	Description	Step summary		
		Time Elapsed	Memory Used This Step	Total Memory Used After
0	Preserving Webpack build output	1767ms	3.7mB	33.29mB
1	Create React pages	29ms	0.17mB	33.48mB
2	Get Drupal content	11518ms	333.17mB	366.67mB
3	Add Drupal Prefix	19ms	0.82mB	367.52mB
4	Create Outreach Assets Data	8ms	0.55mB	368.08mB
5	Create "Resources and support" section of the website	104ms	29.86mB	397.97mB
6	Create environment filter	1ms	0.13mB	398.12mB
7	Add filenames for debugging	1ms	0.39mB	398.53mB
8	Check collections	5ms	0.23mB	398.77mB
9	Group collections	860ms	4.36mB	403.15mB
10	Reset left rail navigation menu levels	3ms	0.54mB	403.7mB
11	Add the date to filenames	29ms	7.28mB	411mB
12	Add app assets	95ms	-3.82mB	407.2mB
13	Add content assets	53ms	-34.55mB	372.66mB
14	Plug the content into the templates	116ms	5.15mB	377.82mB
15	Translate the markdown to html	18ms	1.39mB	379.23mB
16	Add permalinks and change foo.md to foo/index.html	734ms	-0.5mB	378.74mB
17	Create header and footer	11ms	1.54mB	380.29mB
18	Generate navigation	264ms	6.59mB	386.89mB
19	Apply layouts	24142ms	198.07mB	584.97mB
20	Rewrite VA domains for the buildtype	2455ms	-179.65mB	405.37mB
21	Rewrite Drupal pages	0ms	0.01mB	405.43mB
22	Create Drupal debug page	2ms	0.69mB	406.13mB
23	Download Drupal assets	2ms	0.4mB	406.55mB
24	Read application assets from disk	558ms	-28.38mB	378.19mB
25	Create sitemap	130ms	1.77mB	379.98mB
26	Update robots.txt	0ms	0.01mB	380.01mB
27	Check for CMS URLs	1318ms	40.47mB	420.49mB
28	Parse HTML files	42223ms	399.25mB	819.76mB
Build finished!				
🚀 Done in 91.32s.				

▼ Without these changes

		Step summary		
Step	Description	Time Elapsed	Memory Used This Step	Total Memory
Used After Step				
0	Preserving Webpack build output	1962ms	3.22mB	39.66mB
1	Create React pages	10ms	0.16mB	39.84mB
2	Get Drupal content	11080ms	327.19mB	367.04mB
3	Add Drupal Prefix	9ms	2.77mB	369.84mB
4	Create Outreach Assets Data	7ms	0.55mB	370.4mB
5	Create "Resources and support" section of the website	110ms	4.67mB	375.09mB
6	Create environment filter	1ms	0.13mB	375.24mB
7	Add filenames for debugging	1ms	0.39mB	375.65mB
8	Check collections	5ms	0.23mB	375.9mB
9	Group collections	816ms	9.08mB	384.99mB
10	Reset left rail navigation menu levels	3ms	0.73mB	385.73mB
11	Add the date to filenames	27ms	-5.97mB	379.78mB
12	Add app assets	101ms	-1.69mB	378.11mB
13	Add content assets	102ms	-7.48mB	370.64mB
14	Plug the content into the templates	131ms	5.21mB	375.88mB
15	Translate the markdown to html	20ms	1.39mB	377.29mB
16	Add permalinks and change foo.md to foo/index.html	733ms	2.64mB	379.94mB
17	Create header and footer	13ms	1.54mB	381.49mB
18	Generate navigation	281ms	6.03mB	387.54mB
19	Apply layouts	22272ms	154.41mB	541.96mB
20	Rewrite VA domains for the buildtype	2540ms	-130.1mB	411.92mB
21	Rewrite Drupal pages	0ms	0.01mB	411.98mB
22	Create Drupal debug page	2ms	0.69mB	412.69mB
23	Download Drupal assets	2ms	0.41mB	413.11mB
24	Read application assets from disk	584ms	-34.69mB	378.43mB
25	Create sitemap	101ms	2.6mB	381.05mB
26	Update robots.txt	0ms	0.01mB	381.07mB
27	Check for CMS URLs	1324ms	40.99mB	422.08mB
28	Parse HTML files	51322ms	5193.87mB	5615.95mB
29	Add nonce to script tags	1896ms	101.1mB	5717.08mB
30	Process [data-entry-name] attributes into Webpack asset paths	546ms	9.83mB	5726.92mB
31	Update external links	449ms	17.56mB	5744.49mB
32	Add IDs to subheadings	666ms	6.24mB	5750.75mB
33	Check for broken links	5676ms	504.01mB	6254.77mB
34	Inject axe-core for accessibility	17ms	0.88mB	6255.67mB
35	Save the changes from the modified DOM	5612ms	-4871.23mB	1384.46mB
Build finished!				
🚀 Done in 118.86s.				

Testing done

- ⚠️ Needs a lot more testing! Todos -
- ☒ Make preview server's build align with these changes

☒ The preview server (make sure the Axe check still works)

☒ Team site Webpack assets (make sure the hash-free bundles are still being copied over correctly)

- I confirmed this by checking the outputted generated/ directory of a vagovprod build to confirm proxy-rewrite.js and other TeamSite-required files are there

☒ Error handling - make sure broken links still fail the build

- Would like to double confirm this

☒ Confirm each build step still works

Screenshots

I don't think there is any screenshot to post. I pasted the terminal output above.

Acceptance criteria

- ☒ No more memory issues during build

Definition of done

- ☒ Events are logged appropriately

- ☒ Documentation has been updated, if applicable
- ☐ A link has been provided to the originating GitHub issue (or connected to it via ZenHub)
- ☒ No sensitive information (i.e. PII/credentials/internal URLs/etc.) is captured in logging, hardcoded, or specs



3

Attempt to fix memory issue

9f75e29

va-vfs-bot requested a review from **department-of-veterans-affairs/frontend-review-group** on Jan 11

va-vfs-bot reviewed on Jan 11

[View changes](#)

va-vfs-bot left a comment

ESLint is disabled

vets-website uses ESLint to help enforce code quality. In most situations we would like ESLint to remain enabled.

What you can do

See if the code can be refactored to avoid disabling ESLint, or wait for a VSP review.

src/site/stages/build/plugins/modify-dom/check-broken-links/index.js

Show resolved

src/site/stages/build/plugins/modify-dom/parse-html.js **Outdated**

Show resolved

nckslvn changed the title ~~Attempt to fix memory issue~~ **Attempt to fix memory issue during content-build** on Jan 11

lint

b08e839

va-vfs-bot reviewed on Jan 11

[View changes](#)

va-vfs-bot left a comment

Icon found

Icons can be decorative, but sometimes they are used to convey meaning. If there are any semantics associated with an icon, those semantics should also be conveyed to a screen reader.

What you can do

Review the markup and see if the icon provides information that isn't represented textually, or wait for a VSP review.

src/site/stages/build/plugins/modify-dom/add-id-to-subheadings.js

Show resolved

va-vfs-bot reviewed on Jan 11

View changes

va-vfs-bot left a comment

ESLint is disabled

vets-website uses ESLint to help enforce code quality. In most situations we would like ESLint to remain enabled.

What you can do

See if the code can be refactored to avoid disabling ESLint, or wait for a VSP review.

src/site/stages/build/plugins/modify-dom/add-id-to-subheadings.js

Show resolved

src/site/stages/build/plugins/modify-dom/add-id-to-subheadings.js

Outdated

Show resolved

src/site/stages/build/plugins/modify-dom/add-id-to-subheadings.js

Outdated

Show resolved

src/site/stages/build/plugins/modify-dom/inject-axe-core.js

Show resolved

src/site/stages/build/plugins/modify-dom/update-external-links.js

Show resolved

mpelzsherman approved these changes on Jan 11

View changes

mpelzsherman left a comment

Contributor

I tested this on my machine and it worked great! Thank you @ncksllyn!! I hope we are able to merge this soon!

1

mpelzsherman mentioned this pull request on Jan 11

Expand VAMC tests #15543

Merged

ncksllyn added 4 commits on Jan 11

update the preview script to align with changes

50a659c

remove unused file

24e3120

rm unused lint disables

519599f

rename plugin from parseHtml

0a32515

va-vfs-bot reviewed on Jan 11

View changes

va-vfs-bot left a comment

ESLint is disabled

vets-website uses ESLint to help enforce code quality. In most situations we would like ESLint to remain enabled.

What you can do

See if the code can be refactored to avoid disabling ESLint, or wait for a VSP review.

src/site/stages/build/plugins/modify-dom/index.js

Show resolved

nckslvn reviewed on Jan 11

View changes

src/site/stages/build/plugins/modify-dom/index.js

Outdated

```
44 + }
45 + };
46 +
47 + module.exports = parseHtml;
```



nckslvn on Jan 11

Author

Contributor

This whole plugin is the key change. Everything listed in the domModifiers list use to be a separate Metalsmith plugin. This PR converts them into sub-steps within this plugin where each sub-step instead operates on a single file at a time, so that only one file.dom exists in memory at once.

nckslvn marked this pull request as ready for review on Jan 11

nckslvn changed the title Attempt to fix memory issue during content-build Reduce memory consumption during content-build on Jan 11

va-vfs-bot had a problem deploying to master/nps--fix-memory-issue on Jan 11 Error

va-vfs-bot temporarily deployed to master/nps--fix-memory-issue on Jan 11 Inactive

meganhkelley commented on Jan 12

Hey @nckslvn I notice this box still unchecked above: " Team site Webpack assets (make sure the hash-free bundles are still being copied over correctly)" — is that something you are planning to work on, or are you dependent on another team to complete that check?

nckslvn commented on Jan 12

Contributor

Author

Good question @meganhkelley - I'm going to confirm that this morning then I'll check the box. It isn't much work.



1

Move copyAssetsToTeamSitePaths into the conclude hook

f56d50b

nckslvn reviewed on Jan 12

[View changes](#)

src/site/stages/build/plugins/modify-dom/process-entry-names.js

112	110	if (!buildOptions.isPreviewServer) {
113	-	copyAssetsToTeamSitePaths(buildOptions, files, entryNamesDictionary);
	111	+ copyAssetsToTeamSitePaths(buildOptions, files, <u>this</u> .entryNamesDictionary);



nckslvn on Jan 12

Author

Contributor

I'm not certain this `copyAssetsToTeamSitePaths` function really does anything anymore. It seems like whatever is going on here has replaced it,

[vets-website/config/webpack.config.js](#)

Line 25 in b46b14f

```
25 // TODO: refactor the other approach for creating files without the hash so that we're only doing that in the
    webpack config: https://github.com/department-of-veterans-affairs/vets-
    website/blob/a012bad17e5bf024b0ea7326a72ae6a737e349ec/src/site/stages/build/plugins/process-entry-
    names.js#L35
```



nckslvn on Jan 12

Author

Contributor

cc @timwright12 @cvalarida



nckslvn commented on Jan 12

Contributor

Author

@meganhkelley This is ready for review 🙄



Update/add some comments

✓ a7d0081



nckslvn reviewed on Jan 12

[View changes](#)

src/site/stages/build/plugins/modify-dom/add-id-to-subheadings.js

103	+	}
104	+	}
105	+	},
106	+	};



nckslvn on Jan 12

Author

Contributor

This file is basically the same as its predecessor (the file was moved into the `modify-dom` dir) but now wrapped in the `modifyFile` hook

src/site/stages/build/plugins/modify-dom/check-broken-links/index.js

48	+	console.log(errorOutput);
49	+	}
50	+	},
51	+	};



nckslvn on Jan 12

Author

Contributor

Similar to the `add-id-to-subheadings` this file was just refactored to go into a `modifyFile` hook



nckslvn mentioned this pull request on Jan 12

Improve memory usage during the content build department-of-veterans-affairs/va.gov-team#13821


0 of 3 tasks complete

 nckslvn commented on Jan 12

Contributor

Author

@meganhkelley FYSA that the GH issue this PR addresses is here, [department-of-veterans-affairs/va.gov-team#13821](#).

  va-vfs-bot deployed to `master/nps--fix-memory-issue` on Jan 12

[View deployment](#)


 timwright12 approved these changes on Jan 13

[View changes](#)



 timwright12 left a comment

Contributor

Tested locally, all sets appear to be running as before and the site seems to render correctly.

 nckslvn merged commit `07edee3` into `master` on Jan 13
5 checks passed

[View details](#)

  nckslvn deleted the `nps--fix-memory-issue` branch on Jan 13

  acrollet mentioned this pull request on Jan 13

Update `va-gov/web` (v0.1.1217 => v0.1.1227) department-of-veterans-affairs/va.gov-cms#4004

[Merged](#)

  nckslvn mentioned this pull request 9 days ago

[Jenkinsfile] Decrement Drupal parallelization to 10 if not a content deployment #16275

[Merged](#)

0 of 5 tasks complete